

Data Science amb Python

Sprint 8 - Tasca del test d'hipòtesis

Cristiane de Souza da Silva

Abril 2021

Exercises 1

Grab a sports theme dataset you like and select an attribute from the dataset. Calculate the p-value and say if you reject the null hypothesis by taking a 5% alpha.

```
In [1]: # import libraries needed

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #Load dataset
from datetime import datetime

games = pd.read_csv('athlete_olympics.csv')

games.head()
```

```
Out[2]:
```

	ID	Sex	Age	Height	Weight	NOC	Season	Medal
0	1	M	24.0	180.0	80.0	CHN	Summer	NaN
1	2	M	23.0	170.0	60.0	CHN	Summer	NaN
2	3	M	24.0	NaN	NaN	DEN	Summer	NaN
3	4	M	34.0	NaN	NaN	DEN	Summer	Gold
4	5	F	21.0	185.0	82.0	NED	Winter	NaN

```
In [3]: games.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 271116 entries, 0 to 271115
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ID           271116 non-null  int64
1   Sex          271116 non-null  object
2   Age          261642 non-null  float64
3   Height       210945 non-null  float64
4   Weight       208241 non-null  float64
5   NOC          271116 non-null  object
6   Season       271116 non-null  object
7   Medal        39783 non-null   object
dtypes: float64(3), int64(1), object(4)
memory usage: 16.5+ MB
```

```
In [4]: # Replace the Age, Height and Weight NaN by their median.

games['Age'].fillna(games['Age'].median(), inplace=True)

games['Height'].fillna(games['Height'].median(), inplace=True)

games['Weight'].fillna(games['Weight'].median(), inplace=True)
```

```
In [5]: games.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 271116 entries, 0 to 271115
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ID           271116 non-null  int64
1   Sex          271116 non-null  object
2   Age          271116 non-null  float64
3   Height       271116 non-null  float64
4   Weight       271116 non-null  float64
5   NOC          271116 non-null  object
6   Season       271116 non-null  object
7   Medal        39783 non-null   object
dtypes: float64(3), int64(1), object(4)
memory usage: 16.5+ MB
```

```
In [6]: #Replace 'Medal' column Values with Zeros and 1

games['Medal'] = games['Medal'].fillna(0)
games['Medal'] = games['Medal'].replace(['Gold', 'Bronze', 'Silver'], [1, 1, 1])
games.head()
```

```
Out[6]:
```

	ID	Sex	Age	Height	Weight	NOC	Season	Medal
0	1	M	24.0	180.0	80.0	CHN	Summer	0
1	2	M	23.0	170.0	60.0	CHN	Summer	0
2	3	M	24.0	175.0	70.0	DEN	Summer	0
3	4	M	34.0	175.0	70.0	DEN	Summer	1
4	5	F	21.0	185.0	82.0	NED	Winter	0

In []:

In [7]: *# show summary of Dataframe structure*

```
games.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 271116 entries, 0 to 271115
Data columns (total 8 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    ID      271116 non-null  int64
 1   Sex      271116 non-null  object
 2   Age      271116 non-null  float64
 3  Height   271116 non-null  float64
 4  Weight   271116 non-null  float64
 5   NOC      271116 non-null  object
 6  Season   271116 non-null  object
 7  Medal    271116 non-null  int64
dtypes: float64(3), int64(2), object(3)
memory usage: 16.5+ MB
```

In [8]: `games['Sex'].value_counts()`

```
Out[8]: M    196594
        F     74522
        Name: Sex, dtype: int64
```

In [9]: `games.describe()`

```
Out[9]:
```

	ID	Age	Height	Weight	Medal
count	271116.000000	271116.000000	271116.000000	271116.000000	271116.000000
mean	68248.954396	25.502493	175.26374	70.539500	0.146738
std	39022.286345	6.287361	9.27917	12.578184	0.353845
min	1.000000	10.000000	127.00000	25.000000	0.000000
25%	34643.000000	22.000000	170.00000	63.000000	0.000000
50%	68205.000000	24.000000	175.00000	70.000000	0.000000
75%	102097.250000	28.000000	180.00000	75.000000	0.000000
max	135571.000000	97.000000	226.00000	214.000000	1.000000

In [10]: `games.isnull().sum()`

```
Out[10]: ID      0
        Sex      0
        Age      0
        Height    0
        Weight    0
        NOC       0
        Season    0
        Medal     0
        dtype: int64
```

Information about women

```
In [11]: games_women = (games.loc[games['Sex']=='F']).reset_index(drop=True)
games_women.head()
```

```
Out[11]:
```

	ID	Sex	Age	Height	Weight	NOC	Season	Medal
0	5	F	21.0	185.0	82.0	NED	Winter	0
1	5	F	21.0	185.0	82.0	NED	Winter	0
2	5	F	25.0	185.0	82.0	NED	Winter	0
3	5	F	25.0	185.0	82.0	NED	Winter	0
4	5	F	27.0	185.0	82.0	NED	Winter	0

```
In [12]: games_women['Sex'].value_counts()
```

```
Out[12]: F      74522
Name: Sex, dtype: int64
```

```
In [13]: games_women.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74522 entries, 0 to 74521
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ID           74522 non-null  int64
1   Sex          74522 non-null  object
2   Age          74522 non-null  float64
3   Height       74522 non-null  float64
4   Weight       74522 non-null  float64
5   NOC          74522 non-null  object
6   Season       74522 non-null  object
7   Medal        74522 non-null  int64
dtypes: float64(3), int64(2), object(3)
memory usage: 4.5+ MB
```

```
In [14]: games_women.describe().round(1)
```

Out[14]:

	ID	Age	Height	Weight	Medal
count	74522.0	74522.0	74522.0	74522.0	74522.0
mean	69956.7	23.7	168.5	61.1	0.2
std	38932.6	5.8	8.6	10.1	0.4
min	5.0	11.0	127.0	25.0	0.0
25%	36558.5	20.0	163.0	55.0	0.0
50%	70128.0	23.0	169.0	60.0	0.0
75%	103534.8	27.0	175.0	69.0	0.0
max	135568.0	74.0	213.0	167.0	1.0

In []:

In []:

In [15]:

```
# Histogram for Age , Height and Weight
```


```
fig, (ax1, ax2, ax3) = plt.subplots(figsize=(15, 4), ncols=3)
```

```
sns.histplot(data=games_women, x="Age", kde=True, stat="frequency", color='blue')
sns.histplot(data=games_women, x="Height", kde=True, stat="frequency", color='olive')
sns.histplot(data=games_women, x="Weight", kde=True, stat="frequency", color='red')
```

```
fig.suptitle('Female Athletes Information - Histogram', fontsize=16)
```

```
plt.tight_layout()
```

```
plt.show()
```



Next, it will be collected a simple random subset of women dataset of 300 observations.

```
In [16]: #Random Sample of games_women

import random
random.seed(42)

k = 300
games_women_sample = games_women.sample(k).reset_index(drop=True)
games_women_sample.head()
```

```
Out[16]:
```

	ID	Sex	Age	Height	Weight	NOC	Season	Medal
0	116730	F	33.0	172.0	78.0	SUI	Winter	0
1	4809	F	17.0	163.0	59.0	QAT	Summer	0
2	43103	F	24.0	170.0	62.0	RUS	Summer	0
3	71236	F	49.0	164.0	48.0	FRA	Summer	0
4	122703	F	19.0	158.0	48.0	ROU	Summer	0

```
In [17]: games_women_sample.shape
```

```
Out[17]: (300, 8)
```

Weight attribute was selected

Women from olympic games(games_women dataset) weighted 61.1kg on average. I want to determine if the games_women_sample dataset provides strong evidence that women in olympic games weight more or less than the games_women dataset, versus the other possibility that there has been no change.

- **H0:**The average weight of women in the olympic games is 61.1 kg.
- **H1:**The average weight of women in the olympic games is more than 61.1 kg.

The test statistic is the sample mean because $n = 300 > 30$

```
In [18]: games_women_sample.describe().round(1)
```

Out[18]:

	ID	Age	Height	Weight	Medal
count	300.0	300.0	300.0	300.0	300.0
mean	71382.0	24.1	168.7	61.5	0.2
std	38757.0	6.0	8.4	10.5	0.4
min	2912.0	12.0	140.0	33.0	0.0
25%	35974.2	20.0	163.0	55.0	0.0
50%	71536.0	23.0	168.0	61.0	0.0
75%	108036.2	28.0	175.0	68.0	0.0
max	133592.0	49.0	202.0	136.0	1.0

- \bar{x} = population mean
- μ = sample mean
- σ = population standard deviation
- α = significance level
- n = sample size

The critical region could be obtained by selecting a k from the sample mean, so that $CR = \{\mu \leq k\}$ where k is such that $P(\mu \leq k | H_0: \mu = \mu_0) = \alpha = 0.05$. That is under H_0

Follow below the z_score table



$$z \leq (k - \bar{x}) / \sigma / \sqrt{n} = \alpha$$

According to the z_table , for a $\alpha = 0.05$, the z_table is 1.64

```
In [19]: x = 61.1
alpha = 0.05
z_table = 1.64

k = x - z_table * sigma / sqrt(n)
k
```

Out[19]: 59.46

The sample mean is bigger than k ($59.8 > 58.3$), there is no strong evidence to reject H_0 .

Alternative Method

```
In [20]: # Function to find z-value

def z_value(x, mu, sigma, alpha,n):

    '''
    x = population mean
    mu = sample mean
    sigma = population standard deviation
    alpha = significance level
    n = sample size

    '''

    z_value = (mu - x)/(sigma/np.sqrt(n))

    return round(z_value,3)
```

```
In [21]: 10.1/(np.sqrt(300))
```

```
Out[21]: 0.583123771881522
```

```
In [22]: x_z_score = z_value(61.1, 61.6, 10.1,0.05, 300)
x_z_score
```

```
Out[22]: 0.857
```

```
In [23]: from scipy.integrate import quad

def normalProbabilityDensity(x):
    constant = 1.0 / np.sqrt(2*np.pi)
    return(constant * np.exp((-x**2) / 2.0) )

x_z_score_percentile, _ = quad(normalProbabilityDensity, np.NINF, x_z_score)
print('Point Estimate Z-score Probability: ', round(x_z_score_percentile, 3))

Point Estimate Z-score Probability: 0.804
```

```
In [24]: p_value = (1-x_z_score_percentile)
print('p-value = ', round(p_value, 5))

p-value = 0.19572
```

Conclusion

```
In [25]: alpha
print('p-value ', round(p_value, 4), '>', round(alpha, 4), 'alpha value')

p-value 0.1957 > 0.05 alpha value
```


CONCLUSIONS A large p-value indicates ****weakly evidence against**** the null hypothesis H_0 . We ****cannot reject**** the Null Hypothesis that the sample mean is equal to 61.1, at significance level equals to 0.05.

In []:

Exercises 2

Continue with the sports theme dataset you like and select two attributes from the dataset. Calculate the p-value and say if you reject the null hypothesis by taking a 5% alpha.

Set up a hypothesis if there is a relationship between a weight female average and gain a medal

H_0 : There is no difference if a woman gain a medal based on her weight

H_1 : There is some difference if a woman gain a medal based on her weight

```
In [26]: # Separate in group

women_medal = games_women_sample.groupby(['Medal'])
```

```
In [27]: women_medal.Weight.count()
```

```
Out[27]: Medal
0      252
1       48
Name: Weight, dtype: int64
```

```
In [28]: fig = plt.figure(figsize=(15,5))

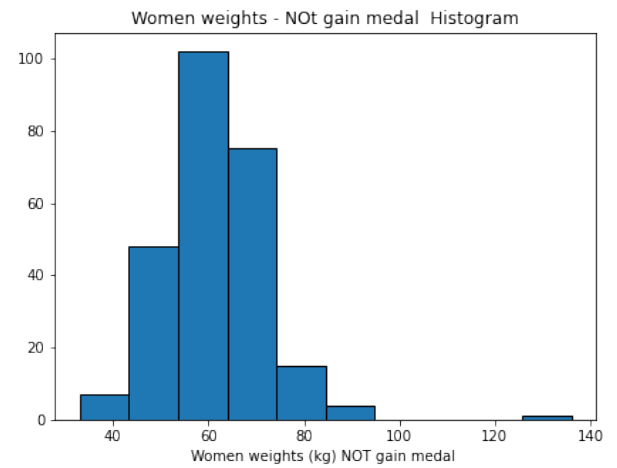
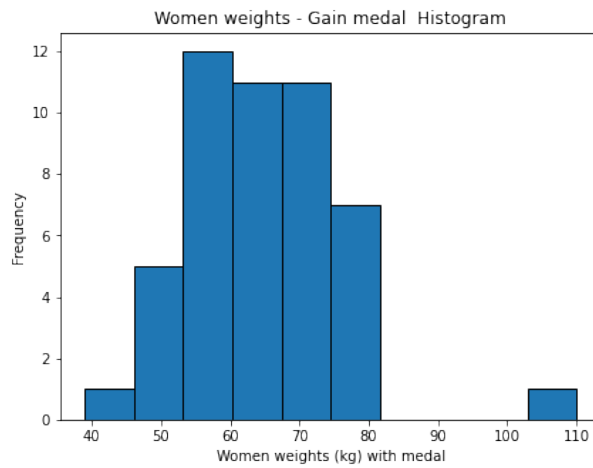
ax1 = fig.add_subplot(1, 2, 1)
s_medal = games_women_sample.loc[games_women_sample['Medal'] == 1]

ax1.hist(s_medal['Weight'], edgecolor='k')
plt.title("Women weights - Gain medal Histogram", fontsize=12)
plt.xlabel("Women weights (kg) with medal")
plt.ylabel("Frequency")

ax2 = fig.add_subplot(1, 2, 2)
s_n_medal = games_women_sample.loc[games_women_sample['Medal'] == 0]

ax2.hist(s_n_medal['Weight'], edgecolor='k')
plt.title("Women weights - NOT gain medal Histogram", fontsize=12)
plt.xlabel("Women weights (kg) NOT gain medal ")

plt.show()
```



```
In [29]: women_medal.mean().round(2)
```

```
Out[29]:
```

	ID	Age	Height	Weight
Medal				

0	72243.97	24.01	168.23	60.97
1	66856.58	24.81	171.08	64.19

```
In [30]: #group_smoke.loc(['nonsmoker','weight'])

x_no_medal = games_women_sample.loc[games_women_sample['Medal'] == 0].Weight
x_medal = games_women_sample.loc[games_women_sample['Medal'] == 1].Weight

x_point_estimate = (x_medal - x_no_medal )
round(x_point_estimate,3)
```

```
Out[30]: 3.217
```

```
In [31]: #Compute the standard error of the point estimate of the population difference

women_medal.std().round(2)
```

```
Out[31]:
```

	ID	Age	Height	Weight
Medal				

0	38777.04	6.07	8.35	10.34
1	38741.01	5.30	8.36	11.21

```
In [32]: se_medal = np.sqrt((10.68**2)/247 + (10.01**2)/53)
round(se_medal,2)
```

```
Out[32]: 1.53
```

```
In [33]: z_score_medal = (x_point_estimate-0)/se_medal
z_score_medal
```

Out[33]: 2.097657669430991

```
In [34]: from scipy.integrate import quad

def normalProbabilityDensity(x):
    constant = 1.0 / np.sqrt(2*np.pi)
    return(constant * np.exp((-x_point_estimate**2) / 2.0) )

x_z_score_percentile, _ = quad(normalProbabilityDensity, np.NINF, z_score_
print('Point Estimate Z-score Probability: ', round(x_z_score_percentile, 4))

Point Estimate Z-score Probability: -0.0023
<ipython-input-34-c52ee754fa06>:9: IntegrationWarning: The integral is prob
ably divergent, or slowly convergent.
    x_z_score_percentile, _ = quad(normalProbabilityDensity, np.NINF, z_score
_medal)
```

```
In [35]: p_value = (1-x_z_score_percentile)
print('p-value = ', round(p_value, 5))
```

p-value = 1.00226

CONCLUSIONS A large p-value indicates **weakly evidence against** the null hypothesis H_0 . We **cannot reject** the Null Hypothesis that there is no difference about Weight mean between who gains and who doesn't gain medal, at significance level equals to 0.05.

Exercises 3

Continue with the sports theme dataset you like and select three attributes from the dataset. Calculate the p-value and say if you reject the null hypothesis by taking a 5% alpha.

```
In [36]: games_women_sample.head()
```

```
Out[36]:
```

	ID	Sex	Age	Height	Weight	NOC	Season	Medal
0	116730	F	33.0	172.0	78.0	SUI	Winter	0
1	4809	F	17.0	163.0	59.0	QAT	Summer	0
2	43103	F	24.0	170.0	62.0	RUS	Summer	0
3	71236	F	49.0	164.0	48.0	FRA	Summer	0
4	122703	F	19.0	158.0	48.0	ROU	Summer	0

```
In [37]: print(round(games_women_sample[['Age', 'Height', 'Weight']].mean(),2))
```

```
Age      24.14
Height   168.68
Weight    61.48
dtype: float64
```

I would like to know if there is real difference between the age of the women according to their country

The three country I'm going to analyse are:

- France (FRA)
- Canada (CAN)
- United States (USA)

```
In [44]: women_country = games_women_sample.loc[(games_women_sample['NOC']=='FRA') |
women_country.head()
```

```
Out[44]:
```

	ID	Sex	Age	Height	Weight	NOC	Season	Medal
3	71236	F	49.0	164.0	48.0	FRA	Summer	0
5	75872	F	29.0	158.0	48.0	FRA	Winter	1
8	69280	F	22.0	178.0	64.0	FRA	Summer	0
19	23739	F	28.0	165.0	63.0	CAN	Summer	0
21	46783	F	21.0	175.0	136.0	USA	Summer	0

```
In [45]: women_country.NOC.value_counts()
```

```
Out[45]: USA      28
CAN       14
FRA       11
Name: NOC, dtype: int64
```

```
In [47]: women_country.groupby('NOC')[['Age']].mean().round(3)
```

```
Out[47]:
```

	Age
NOC	
CAN	24.429
FRA	25.364
USA	26.500

- H0 : the age average of the three countries are the same
- H1 : the age average of the three countries are not the same

```
In [48]: women_country.groupby('NOC')[['Age']].std().transpose().round(3)
```

Out[48]:

	NOC	CAN	FRA	USA
Age	4.767	8.88	7.671	

```
In [49]: import scipy.stats as stats
import researchpy as rp
import statsmodels.api as sm
from statsmodels.formula.api import ols
```

```
In [50]: # isolate the variables of interest

rp.summary_cont(women_country['Age']).round(3)
```

Out[50]:

	Variable	N	Mean	SD	SE	95% Conf. Interval
0	Age	53.0	25.717	7.225	0.992	23.725 27.709

```
In [51]: rp.summary_cont(women_country['Age'].groupby(women_country['NOC'])).round(3)
```

Out[51]:

		N	Mean	SD	SE	95% Conf. Interval
	NOC					
	CAN	14	24.429	4.767	1.274	21.676 27.181
	FRA	11	25.364	8.880	2.677	19.398 31.329
	USA	28	26.500	7.672	1.450	23.525 29.475

```
In [52]: zstats, pvalue = stats.f_oneway(women_country['Age'][women_country['NOC'] == 'CAN'],
                                         women_country['Age'][women_country['NOC'] == 'FRA'],
                                         women_country['Age'][women_country['NOC'] == 'USA'])

print('z-statistics value = ', zstats.round(3), ' p-value = ', pvalue.round(3))

z-statistics value = 0.391 p-value = 0.6786
```

CONCLUSIONS A large p-value indicates ****weakly evidence against**** the null hypothesis H_0 , i.e, the age average of the three (Canada, France and United States) countries are the same. We ****cannot reject**** the Null Hypothesis that there is no difference about Weight mean between who gains and who doesn't gain medal, at significance level equals to 0.05.

In []: