

# Data Science amb Python

## Sprint 12

### S12 T01 : Supervised Learning - Regressions

Cristiane de Souza da Silva

June 2021

## Dataset Description

### Airline2008Nov Dataset Variable definition

Name	Description
1.Year	2008
2.Month	11
3.DayofMonth	1-31
4.DayOfWeek	1 (Monday) - 7 (Sunday)
5.DepTime	actual departure time (local, hhmm)
6.CRSDepTime	scheduled departure time (local, hhmm)
7.ArrTime	actual arrival time (local, hhmm)
8.CRSArrTime	scheduled arrival time (local, hhmm)
9.UniqueCarrier	unique carrier code
10.FlightNum	flight number
11.TailNum	plane tail number
12.ActualElapsedTime	in minutes
13.CRSElapsedTime	CRS Elapsed Time of Flight (estimated elapse time), in minutes
14.AirTime	Flight Time, in Minutes, in minutes
15.ArrDelay	Difference in minutes between scheduled and actual arrival time. Early arrivals show negative numbers, in minutes
16.DepDelay	Difference in minutes between scheduled and actual departure time. Early departures show negative numbers, in minutes
17.Origin	origin IATA airport code
18.Dest	destination IATA airport code
19.Distance	Distance between airports (miles)
20.TaxiIn	Wheels down and arrival at the destination airport gate, in minutes

21.TaxiOut	The time elapsed between departure from the origin airport gate and wheels off, in minutes
22.Cancelled	was the flight cancelled?
23.CancellationCode	reason for cancellation (A = carrier, B = weather, C = NAS, D = security)
24.Diverted	1 = yes, 0 = no
25.CarrierDelay	minutes. Carrier delay is within the control of the air carrier. Examples:: aircraft cleaning, aircraft damage, etc
26.WeatherDelay	munutes. Weather delay is caused by extreme or hazardous weather conditions
27.NASDelay	minutes. Delay that is within the control of the National Airspace System (NAS) Ex:airport operations,heavy traffic volume,etc
28.SecurityDelay	minutes. caused by evacuation of a terminal or concourse, re-boarding of aircraft because of security breach,etc
29.LateAircraftDelay	minutes. Arrival delay at an airport due to the late arrival of the same aircraft at a previous airport.

More information about flights delay can be seen in [Federal Aviation Administration site](#).

## Exercise 1

Split the DelayedFlights.csv dataset into train and test. Study the two sets separately, at a descriptive level

### 1. Load the dataset

```
In [1]: # import the needed libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # Load the dataframe

df_flights = pd.read_csv('DelayedFlights.csv')
df_flights.head()
```

```
Out[2]:
```

	Unnamed: 0	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CR
0	0	2008	1	3	4	2003.0	1955	2211.0	
1	1	2008	1	3	4	754.0	735	1002.0	
2	2	2008	1	3	4	628.0	620	804.0	
3	4	2008	1	3	4	1829.0	1755	1959.0	
4	5	2008	1	3	4	1940.0	1915	2121.0	

5 rows × 30 columns

```
In [3]: df_flights.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1936758 entries, 0 to 1936757
Data columns (total 30 columns):
#   Column                Dtype
---  -
0   Unnamed: 0            int64
1   Year                  int64
2   Month                 int64
3   DayofMonth            int64
4   DayOfWeek             int64
5   DepTime               float64
6   CRSDepTime            int64
7   ArrTime               float64
8   CRSArrTime            int64
9   UniqueCarrier         object
10  FlightNum             int64
11  TailNum               object
12  ActualElapsedTime     float64
13  CRSElapsedTime        float64
14  AirTime               float64
15  ArrDelay              float64
16  DepDelay              float64
17  Origin                object
18  Dest                  object
19  Distance              int64
20  TaxiIn                float64
21  TaxiOut               float64
22  Cancelled             int64
23  CancellationCode      object
24  Diverted              int64
25  CarrierDelay          float64
26  WeatherDelay          float64
27  NASDelay              float64
28  SecurityDelay         float64
29  LateAircraftDelay     float64
dtypes: float64(14), int64(11), object(5)
memory usage: 443.3+ MB
```

## 2. Data Pre-processing

```
In [4]: #Remove the column 'Unnamed: 0'
df_flights = df_flights.drop('Unnamed: 0', axis=1)
df_flights.head()
```

```
Out[4]:
```

	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	...
0	2008	1	3	4	2003.0	1955	2211.0	2225	...
1	2008	1	3	4	754.0	735	1002.0	1000	...
2	2008	1	3	4	628.0	620	804.0	750	...
3	2008	1	3	4	1829.0	1755	1959.0	1925	...
4	2008	1	3	4	1940.0	1915	2121.0	2110	...

5 rows × 29 columns

## Data Cleaning

```
In [5]: #Verify missing data
df_flights.isnull().sum()
```

```
Out[5]: Year                0
Month                  0
DayofMonth             0
DayOfWeek              0
DepTime                0
CRSDepTime             0
ArrTime                7110
CRSArrTime             0
UniqueCarrier          0
FlightNum              0
TailNum                5
ActualElapsedTime      8387
CRSElapsedTime         198
AirTime               8387
ArrDelay              8387
DepDelay               0
Origin                0
Dest                  0
Distance              0
TaxiIn                7110
TaxiOut               455
Cancelled              0
CancellationCode       0
Diverted               0
CarrierDelay          689270
WeatherDelay          689270
NASDelay              689270
SecurityDelay          689270
LateAircraftDelay     689270
dtype: int64
```

```
In [6]: # copy the data frame df_flights

df_clean = df_flights.copy()
```

I'll drop some columns that has no impact with the objective, such as the Year ( all flights happen in 2008), Cancelled, etc

```
In [7]: df_clean = df_clean.drop(['Year', 'FlightNum', 'TailNum', 'CRSDepTime',
                                'CRSArrTime', 'TaxiIn', 'TaxiOut', 'Cancelled', 'C'
                                ])
df_clean.head()
```

```
Out[7]:
```

	Month	DayofMonth	DayOfWeek	DepTime	ArrTime	UniqueCarrier	ActualElapsedTime
0	1	3	4	2003.0	2211.0	WN	128.0
1	1	3	4	754.0	1002.0	WN	128.0
2	1	3	4	628.0	804.0	WN	96.0
3	1	3	4	1829.0	1959.0	WN	90.0
4	1	3	4	1940.0	2121.0	WN	101.0

## Analysing each feature with missing data

- **ArrDelay** : Delay of arrival time is what we want to predict.

Any data entry without ArrDelay data cannot be used for prediction; therefore any row now without ArrDelay data is not useful to us.

There are 8387 missing values in ArrDelay. First, let's drop this rows.

```
In [8]: # drop whole row with NaN in ArrDelay column

df_clean.dropna(subset = ['ArrDelay'], axis=0, inplace=True )

# reset index
df_clean.reset_index(drop=True, inplace=True)
```

```
In [9]: # Verify again the missing data

df_clean.isnull().sum()
```

```
Out[9]: Month                0
        DayofMonth           0
        DayOfWeek            0
        DepTime              0
        ArrTime              0
        UniqueCarrier        0
        ActualElapsedTime     0
        CRSElapsedTime       0
        AirTime              0
        ArrDelay              0
        DepDelay              0
        Origin                0
        Dest                  0
        Distance              0
        CarrierDelay          680883
        WeatherDelay          680883
        NASDelay              680883
        SecurityDelay         680883
        LateAircraftDelay     680883
        dtype: int64
```

We can see that most of the missing values of the other attributes were related to ArrDelay

For the other attributes, **CarrierDelay**, **WeatherDelay**, **NASDelay**, **SecurityDelay**, **LateAircraftDelay**, I'll verify the percentage of this missing values.

```
In [10]: missing_percentage = (df_clean['CarrierDelay'].isnull().sum())/(df_clean['CarrierDelay'].count())
        missing_percentage
```

```
Out[10]: 0.35308713935233416
```

35.3% of these values are missing and it is a big percentage. Because of that, I'll drop them.

```
In [11]: df_clean = df_clean.drop(['CarrierDelay', 'WeatherDelay', 'NASDelay', 'SecurityDelay'])
        df_clean.head()
```

```
Out[11]:
```

	Month	DayofMonth	DayOfWeek	DepTime	ArrTime	UniqueCarrier	ActualElapsedTime
0	1	3	4	2003.0	2211.0	WN	128.0
1	1	3	4	754.0	1002.0	WN	128.0
2	1	3	4	628.0	804.0	WN	96.0
3	1	3	4	1829.0	1959.0	WN	90.0
4	1	3	4	1940.0	2121.0	WN	101.0

## Data Transforming

Encode the categorical data in number.

```
In [12]: print('The number of unique values of UniqueCarrier are: ', df_clean['UniqueCarrier'].nunique())
print('The number of unique values of Origin are: ', df_clean['Origin'].nunique())
print('The number of unique values of Dest are: ', df_clean['Dest'].nunique())
```

```
The number of unique values of UniqueCarrier are: 20
The number of unique values of Origin are: 303
The number of unique values of Dest are: 302
```

```
In [18]: from sklearn import preprocessing
le = preprocessing.LabelEncoder()

df_cat = df_clean[['UniqueCarrier', 'Origin', 'Dest']]
df_cat.head()
```

```
Out[18]:
```

	UniqueCarrier	Origin	Dest
0	WN	IAD	TPA
1	WN	IAD	TPA
2	WN	IND	BWI
3	WN	IND	BWI
4	WN	IND	JAX

```
In [19]: #for col in df_cat:
#df_clean[col] = le.fit_transform(df_cat[col])
```

```
In [20]: #df_clean.head()
```

```
In [21]: #save the df_clean

df_clean.to_csv('delayedflights_clean.csv', index=False)
```

## Training and Test

```
In [22]: from sklearn.model_selection import train_test_split

y = df_clean['ArrDelay']

X = df_clean.drop('ArrDelay', axis=1)
```

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)

print("number of test samples :", X_test.shape[0])
print("number of training samples:", X_train.shape[0])
```

```
number of test samples : 578512
number of training samples: 1349859
```

## Exercises 2

Apply some transformation process (standardize numerical data, create dummy columns, polynomials ...).

```
In [24]: #Encode the categorical data in number
from sklearn import preprocessing
le = preprocessing.LabelEncoder()

X_train_cat = X_train[['UniqueCarrier','Origin', 'Dest' ]]
X_test_cat = X_test[['UniqueCarrier','Origin', 'Dest' ]]
#df_cat.head()
```

```
In [25]: for col in X_train_cat:
        X_train[col] = le.fit_transform(X_train_cat[col])

        for col in X_test_cat:
            X_test[col] = le.fit_transform(X_test_cat[col])

        X_train.head()
```

<ipython-input-25-3ba92260ebc5>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
X_train[col] = le.fit_transform(X_train_cat[col])
<ipython-input-25-3ba92260ebc5>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
X_test[col] = le.fit_transform(X_test_cat[col])
```

```
Out[25]:
```

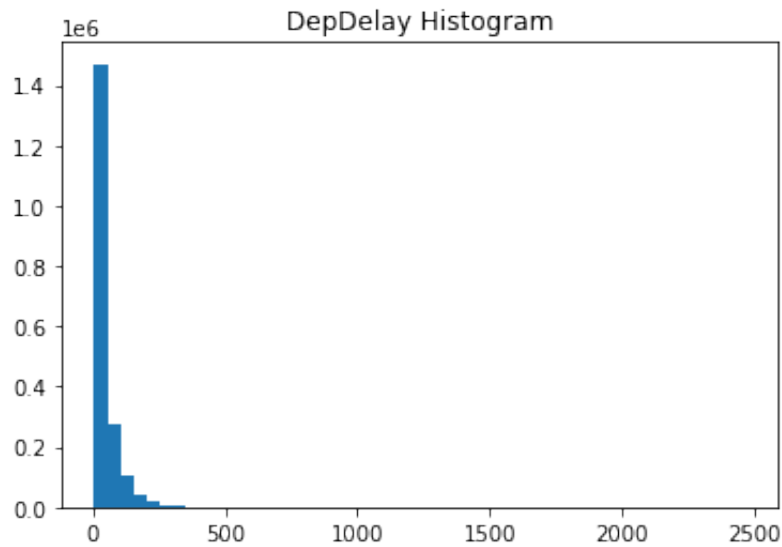
	Month	DayofMonth	DayOfWeek	DepTime	ArrTime	UniqueCarrier	ActualElapse
995566	6	4	3	2308.0	113.0	9	
1090173	7	14	1	2127.0	2211.0	17	
1550228	10	24	5	1330.0	1522.0	13	
870468	5	7	3	1122.0	1616.0	5	
383192	3	13	4	1827.0	1923.0	17	

## Normalize Data

```
In [26]: import matplotlib.pyplot as plt
import seaborn as sns

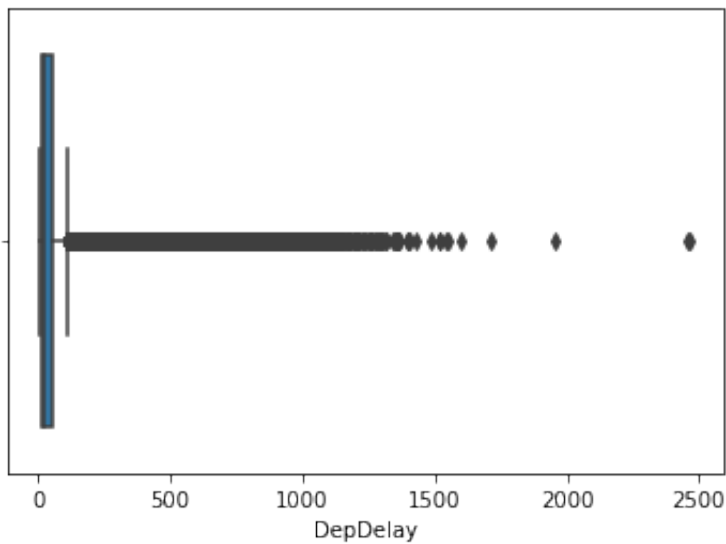
plt.hist(X["DepDelay"], bins=50)
plt.title('DepDelay Histogram');
```





```
In [27]: sns.boxplot(data = X, x="DepDelay")
```

```
Out[27]: <AxesSubplot:xlabel='DepDelay'>
```



The plots above shows that the data contains many outliers, so I'll use you the RobustScaler

```
In [28]: #from sklearn.preprocessing import StandardScaler
#scaler = StandardScaler().fit(X_train)

#standardized_X = scaler.transform(X_train)
#standardized_X_test = scaler.transform(X_test)
```

```
In [29]: from sklearn.preprocessing import RobustScaler

transformer = RobustScaler().fit(X_train)
transformer
```

```
Out[29]: RobustScaler()
```

```
In [30]: standardized_X = transformer.transform(X_train)
         standardized_X_test = transformer.transform(X_test)
```

## Exercises 3

Summarize the new columns generated statistically and graphically

```
In [32]: X_train_stand=pd.DataFrame(X_train, columns=X_train.columns).reset_index(drop=True)
         X_train_stand
```

```
Out[32]:
```

	Month	DayofMonth	DayOfWeek	DepTime	ArrTime	UniqueCarrier	ActualElapse
0	6	4	3	2308.0	113.0	9	
1	7	14	1	2127.0	2211.0	17	
2	10	24	5	1330.0	1522.0	13	
3	5	7	3	1122.0	1616.0	5	
4	3	13	4	1827.0	1923.0	17	
...	...	...	...	...	...	...	
1349854	3	10	1	1017.0	1230.0	9	
1349855	12	21	7	1405.0	1648.0	7	
1349856	3	18	2	1001.0	1042.0	6	
1349857	3	13	4	1853.0	1925.0	9	
1349858	1	28	1	1818.0	2007.0	11	

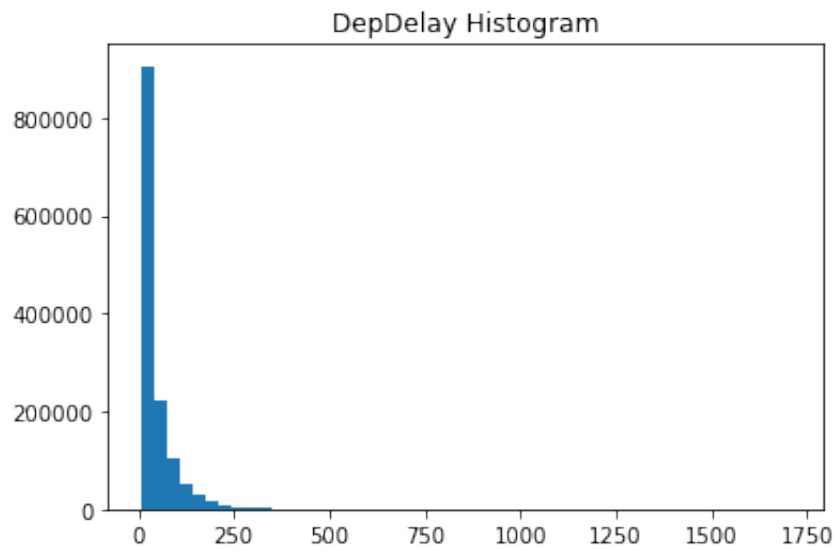
1349859 rows × 13 columns

```
In [33]: X_train_stand.describe()
```

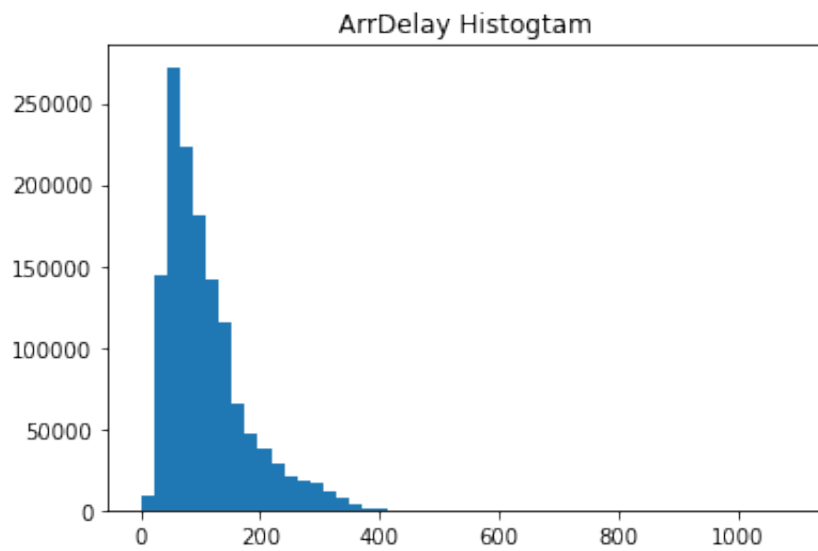
```
Out[33]:
```

	Month	DayofMonth	DayOfWeek	DepTime	ArrTime	UniqueC
count	1.349859e+06	1.349859e+06	1.349859e+06	1.349859e+06	1.349859e+06	1.349859e+06
mean	6.110603e+00	1.574881e+01	3.985953e+00	1.518495e+03	1.610178e+03	1.112838e+01
std	3.482405e+00	8.776654e+00	1.995536e+00	4.504583e+02	5.478223e+02	5.932061e+00
min	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	0.000000e+00
25%	3.000000e+00	8.000000e+00	2.000000e+00	1.203000e+03	1.316000e+03	6.000000e+00
50%	6.000000e+00	1.600000e+01	4.000000e+00	1.545000e+03	1.715000e+03	1.300000e+01
75%	9.000000e+00	2.300000e+01	6.000000e+00	1.900000e+03	2.031000e+03	1.700000e+01
max	1.200000e+01	3.100000e+01	7.000000e+00	2.400000e+03	2.400000e+03	1.900000e+01

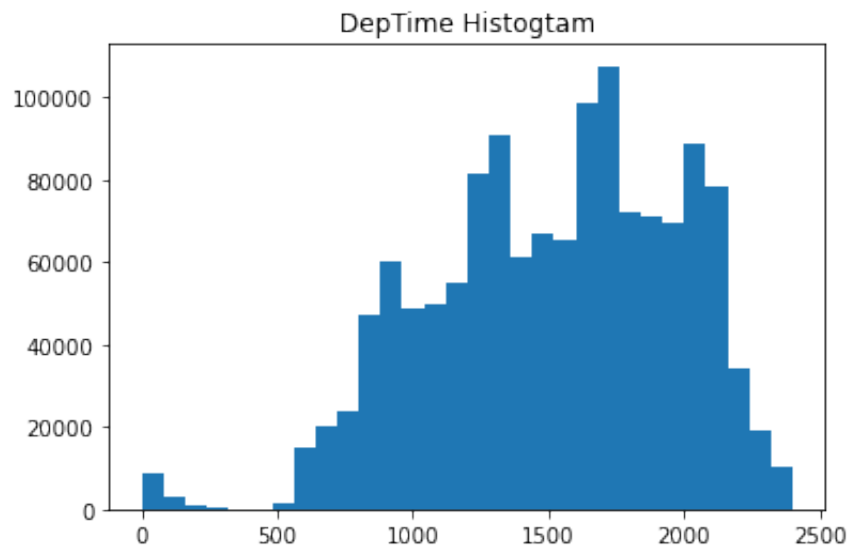
```
In [34]: plt.hist(X_train_stand["DepDelay"], bins=50)
         plt.title('DepDelay Histogram');
```



```
In [35]: plt.hist(X_train_stand["AirTime"], bins=50)
plt.title('ArrDelay Histogtam');
```



```
In [36]: plt.hist(X_train_stand["DepTime"], bins=30)
plt.title('DepTime Histogtam');
```



In [ ]: